

Access to Big Data in Bioinformatics

Brendan Ball
Andrew van Rooyen

1. PROJECT DESCRIPTION

Collaboration between bioinformatics organizations involves shared access to large datasets. This project investigates two avenues for tackling the collaborative analysis of big data in bioinformatics: Efficient transport of large datasets across high speed WAN networks and implementation of community clouds that host and securely execute code close to the location where data is stored.

Historically, data has been exchanged using tools and protocols like SSH and FTP, but new protocols, such as GridFTP [Allcock et al. 2005] and HPN-SSH [Rapier and Bennett 2008], offer more efficient use of high speed networks.

Community clouds [Briscoe and Marinos 2009] are cloud computing services built on micro clouds hosted by collaborating organisations, as opposed to conventional cloud computing hosted by cloud vendors (such as Microsoft, Google and Amazon) [Jadeja and Modi 2012]. Hosting micro clouds close to scientific data collections would facilitate scientific collaboration through moving code closer to data, rather than vice versa.

2. PROBLEM STATEMENT

The University of Cape Town (UCT) and the University of the Western Cape (UWC) are both part of the South African National Research Network (SANReN). This provides a 10Gb link to other academic institutions in South Africa. Currently, bioinformatics data in the Western Cape cannot be moved at these speeds, even though the infrastructure theoretically allows it. In practice, the software configuration at the endpoints is the bottleneck. In the scope of this project, the software configuration refers to the network protocol, but in practice includes more layers. This project aims to bring the transfer rate up to the actual capacity of the network by investigating different network protocols.

The research question here is: Which network protocol is best for use on a 10Gb/s network, and can it overcome the challenges faced by FTP and SCP?

The second part of this project will survey existing cloud computing infrastructures towards their suitability for creating organisation-level micro clouds. The aim is to design and implement a micro cloud solution allowing universities to easily deploy their own micro cloud which can access and be accessed by other connected micro clouds. These connected micro clouds will form a community cloud. The micro cloud solution will include implementing a code migration framework which will allow users to easily migrate code between micro clouds, execute that code on the micro cloud which is storing the data and return any results to a specific micro cloud.

The research question for the second part is as follows: Is it possible to implement a micro cloud platform which provides the convenience of processing bioinformatics data locally, but bypasses the challenges around transferring the data sets?

3. PROCEDURES AND METHODS

The first part of this project will compare and contrast the performance of GridFTP, HPN-SSH, FTP and SSH for transferring multi-gigabyte datasets. This will be eval-

uated by comparing throughput, delays, security and authentication features of these protocols and their implementations. Throughput and delays will be measured quantitatively via operating system tools and logging. Other features will be compared by examining literature (whitepapers, definitions) on the protocols themselves.

This will be tested mostly on the local UCT network, and then between UCT and UWC once the nuances of the protocols and their configurations are understood.

For the second part, existing software will be surveyed, and a combination of solutions at various stack levels will be fitted together to form a complete solution. At least two micro clouds will be deployed in order to prototype a community cloud where code can be migrated between micro cloud installations. The community cloud will be evaluated on both its functionality and usability, with more emphasis on functionality. The functionality will be evaluated by testing for the successful processing of given bioinformatics analysis code on example bioinformatics data. Usability testing will be done by doing systematic observation. A limited number of experts will be observed to determine usability using client satisfaction measures.

4. ETHICAL, PROFESSIONAL AND LEGAL ISSUES

Usually, there would be ethical issues when working with bioinformatics data. However, because this project is only a platform, testing doesn't rely on having data specific to a particular person or group, and it doesn't have to be recent. Any data we need can be obtained from public sources such as the National Center for Biotechnology Information (NCBI) database.

All experiments for the first section will involve quantitative testing of systems. No ethical issues are raised here. However, access to the UWC network (especially via SANReN) will need to be granted by the appropriate authorities. There may be user testing involved with the second part, but this will be expert testing with the collaborators of the project. We conclude that no ethical clearance is necessary for this project.

All work will be made publically available under the MIT licence.

5. RELATED WORK

Data sets are a big part of bioinformatics, and have introduced many new challenges with the rise of next generation sequencing. Sequencing technologies like SOLiD provide much higher data output at a cheaper cost [Shendure and Ji 2008], which is good news for research, but troubling for data storage, transfer and access. In fact, the cost of storing a byte has been more expensive than sequencing a base pair since before 2010 [Baker 2010].

This makes it difficult for researchers in different locations to manipulate and run processes on the data, because it will be stored in only one location. These files could be tens of gigabytes in size [Deorowicz and Grabowski 2011], depending on context.

There has been a lot of work on storing this data. There are a plethora of file formats whose efficiency depends on the kind of data which needs to be stored. Two of the most popular formats are FASTQ, which stores aggregated reads along with the quality of each base pair [Cock et al. 2010], and BAM, the binary, compressed version of the Sequence Alignment Map (SAM) format [SAMTools 2015].

There are also some proprietary transfer protocols which are widely used in practice. For example, the fasp protocol by the US based company AsperaSoft. Based on UDP, the protocol eliminates the latency issues seen with TCP, and provides bandwidth up to 10 gigabits per second to transfer data [Beloslyudtsev 2014].

There has been an explorative push towards cloud solutions from Amazon, Google etc, but there are very significant drawbacks. Because the sequencing happens in labs, researchers need to upload their raw results to the cloud data centres every time they run a new experiment. This leads back to the original problem, as researchers resort to mailing hard drives [Baker 2010].

There are also security, privacy and ethical concerns with outsourcing this processing power to other companies, as sequenced DNA data is often highly sensitive information [Marx 2013].

Work in this area includes Cloud BioLinux, which is a community driven project focussed on next generation sequencing. It is a toolkit which makes it easy to deploy virtual machines with bioinformatics infrastructure to a cloud platform. It bundles specific packages used in next generation sequence analysis, thereby decreasing configuration time and increasing maintainability. Instances of Cloud BioLinux have been tested on the Amazon EC2 cloud platform and on a private Eucalyptus cloud. [Krampis et al. 2012]

6. POSSIBLE FUTURE WORK

- Further optimising the data transfer pipeline on different levels of the stack. For example, tuning how data is read from disk to match patterns used by the network layer.
- Linking user identity to existing databases, and using this as a platform for access control. Users can, for example, have permission to execute code on remote micro clouds if they are trusted.
- Visualising genomic datasets with software such as Visual Molecular Dynamics (VMD), and linking this to the traditional linear view of the data. Possible advances could involve highlighting areas on a 3D model when a section of the linear data is interacted with. This process would need to run on a cloud and send relevant data back to the user.

7. ANTICIPATED OUTCOMES

- Choice of best transfer protocol for big data in our context of bioinformatics in South Africa.
- A micro cloud solution providing capability of creating a community cloud.

7.1. Impact

- UWC can use their dedicated 10Gbs line
- Collaboration between universities can increase given the new community cloud platform, particularly in the bioinformatics department.

8. PROJECT PLAN

8.1. Resources Required

- Access to servers micro cloud solution
- Access to SANReN for testing transfer protocols

8.2. Deliverables

- 24 July - Initial Feasibility Demonstration
- 30 July - Background of final paper
- 11 September - First Implementation/Experiment/Performance test + Writeup
- 21 September - Final Prototype/Experiment /Performance test + Writeup
- 25 September - Final Implementation and Testing
- 2 October - Outline of Complete Report
- 16 October - Final Complete Draft of Report
- 2 November - Project Report Final Submission
- 2 November - Poster
- 9 November - Website
- 18 November - Reflection Paper

8.3. Milestones

See see Fig 1 for Gantt chart.

- 12 June - Read up on GridFTP and HPN-SSH
- 12 June - Decide what a micro cloud agent would look like
- 13 July - Dry run on UCT network
- 20 July - Implement a management and deployment system for the agents
- 24 July - Initial Feasibility Demonstration
- 30 July - Background of final paper
- 28 August - Design Section
- 11 September - Authentication and security worked into all moving parts
- 11 September - First Implementation/Experiment/Performance test + Writeup
- 21 September - Final Prototype/Experiment /Performance test + Writeup
- 25 September - Final Implementation and Testing
- 2 October - Outline of Complete Report
- 16 October - Final Complete Draft of Report
- 2 November - Project Report Final Submission
- 2 November - Poster
- 9 November - Website
- 18 November - Reflection Paper

8.4. Work Allocation

- The first section (comparing transfer protocols) will be completed by van Rooyen. Since this section is anticipated to be shorter than Balls, van Rooyen will then work on security and authentication for the cloud platform
- Ball will do the second section (micro cloud solution)




Fig. 2: Work split
Blue: Ball
Red: van Rooyen

This distribution is not due to any particular skills of van Rooyen or Ball.

8.5. Risks

Risk	Mitigation
One of the team members is unable to complete his work on time due to unforeseen circumstances	The two sections of work are loosely coupled so it won't prevent the other team member from completing his work
Not getting access to key resources	Keep open communication with supervisors
Project specification is too big, team is unable to complete all tasks on time.	Change project specification based on feedback from the presentation
Team conflict	The project is loosely coupled which should decrease chance of team conflict
Failure to integrate project components	Discuss design decisions to account for integration

REFERENCES

- William Allcock, John Bresnahan, Rajkumar Kettimuthu, Michael Link, Catalin Dumitrescu, Ioan Raicu, and Ian Foster. 2005. The Globus striped GridFTP framework and server. In *Proceedings of the 2005 ACM/IEEE conference on Supercomputing*. IEEE Computer Society, 54.
- Monya Baker. 2010. Next-generation sequencing: adjusting to data overload. *nature methods* 7, 7 (2010), 495–499.
- Dima Beloslyudtsev. 2014. Aspera transfer guide. (2014).
- Gerard Briscoe and Alexandros Marinos. 2009. Digital ecosystems in the clouds: towards community cloud computing. In *Digital Ecosystems and Technologies, 2009. DEST'09. 3rd IEEE International Conference on*. IEEE, 103–108.
- Peter JA Cock, Christopher J Fields, Naohisa Goto, Michael L Heuer, and Peter M Rice. 2010. The Sanger FASTQ file format for sequences with quality scores, and the Solexa/Illumina FASTQ variants. *Nucleic acids research* 38, 6 (2010), 1767–1771.

- Sebastian Deorowicz and Szymon Grabowski. 2011. Compression of DNA sequence reads in FASTQ format. *Bioinformatics* 27, 6 (2011), 860–862.
- Yaju Jadeja and Kavan Modi. 2012. Cloud computing-concepts, architecture and challenges. In *Computing, Electronics and Electrical Technologies (ICCEET), 2012 International Conference on*. IEEE, 877–880.
- Konstantinos Krampis, Tim Booth, Brad Chapman, Bela Tiwari, Mesude Bicak, Dawn Field, and Karen E Nelson. 2012. Cloud BioLinux: pre-configured and on-demand bioinformatics computing for the genomics community. *BMC bioinformatics* 13, 1 (2012), 42.
- Vivien Marx. 2013. Biology: The big challenges of big data. *Nature* 498, 7453 (2013), 255–260.
- Chris Rapiet and Benjamin Bennett. 2008. High speed bulk data transfer using the SSH protocol. In *Proceedings of the 15th ACM Mardi Gras conference: From lightweight mash-ups to lambda grids: Understanding the spectrum of distributed computing requirements, applications, tools, infrastructures, interoperability, and the incremental adoption of key capabilities*. ACM, 11.
- SAMTools. 2015. Sequence Alignment/Map Format Specification. <https://samtools.github.io/hts-specs/SAMv1.pdf>. (2015). Accessed: 2015-04-27.
- Jay Shendure and Hanlee Ji. 2008. Next-generation DNA sequencing. *Nature biotechnology* 26, 10 (2008), 1135–1145.




Fig. 1: Gantt Chart